

CSCI 5409

Cloud Computing

Term Assignment – 2024 Summer

Name: Divy Bhavsar

BannerID: B00978267

Overview

I have developed a comprehensive web-based stock trading application named '**SKEX**'. This application leverages modern web technologies to provide a seamless trading experience. The frontend is crafted using Next.js, ensuring a dynamic and responsive user interface. The backend is built with Node.js 18, offering robust and efficient server-side operations. Data management is handled using AWS RDS with a MariaDB database, ensuring scalable and reliable data storage and retrieval.

Feature

Stock Data

- **Data Source:** Stock data is fetched from a third-party API, Polygon.io.
- **Initial Data:** Initially, the application displays data for 50 stocks.
- **Public Access:** Unauthenticated users can view stock data but cannot perform any trading operations.
- **Authenticated Operations:** Logged-in users can buy and sell stocks. If an unauthenticated user attempts to perform a trade, they are redirected to the login page.

Authentication

Sign Up

- **Fields Required:** Email, SIN number, and password.
- **User Experience:** The sign-up process is straightforward, collecting essential information to create a secure user account.

Login

- **Fields Required:** Email and password.
- **User Experience:** Users can log in using their email and password, gaining access to trading functionalities.

Trading Operations

Buy and Sell Buttons

- **Location:** Buttons are present in each row of the stock information table.
- **Functionality:** When a user clicks the buy or sell button for a specific stock, they are redirected to the respective form.
- **Forms:**
 - **Buy Form:** Users need to specify the quantity of the stock they wish to purchase and hit the buy button.
 - **Sell Form:** Users need to specify the quantity of the stock they wish to sell and hit the sell button.

- **Validation:**
 - **Buy:** If the user is not logged in, they are redirected to the login page.
 - **Sell:** If the user does not have the stock in their portfolio, the sell request is not accepted.

User Portfolio

- **Access:** The portfolio page can only be accessed by authenticated users.
- **Content:** The portfolio page displays the stocks purchased by the user, providing an overview of their holdings.

User Flow

1. **View Stock Data:** All users can view stock data fetched from Polygon.io.
2. **Authentication:**
 - **Sign Up:** New users provide their email, SIN number, and password to create an account.
 - **Login:** Existing users log in with their email and password.
3. **Trading:**
 - **Buy/Sell Operations:** Authenticated users can buy or sell stocks by clicking the respective button in the stock information table.
 - **Form Redirection:** Upon clicking the buy or sell button, users are redirected to a form where they specify the quantity and confirm the transaction.
4. **Portfolio Management:** Authenticated users can access their portfolio page to view their purchased stocks.

Drawback: As per the API limit , The application can not fetch live price of stocks, which can be solved in future versions.

AWS Services

1. AWS VPC (Virtual Private Cloud)

Description:

AWS VPC was used to create a logically isolated network environment where I could launch AWS resources in a virtual network that I define.

Alternatives Considered:

- **AWS Transit Gateway:** Provides scalable and easy-to-manage connections between VPCs and on-premises networks but is more suited for complex network architectures.
- **AWS Direct Connect:** Offers a dedicated network connection to AWS but is more suitable for hybrid environments requiring high-bandwidth connections.

Why Chosen:

- **Integration:** AWS VPC integrates seamlessly with other AWS services used in the project.
- **Control:** Provides granular control over the network environment, including subnets, route tables, and network gateways.
- **Security:** Enhanced security features, including security groups and network ACLs, ensured a secure environment for sensitive financial data.

2. Amazon EC2 (Elastic Compute Cloud)

Description:

Amazon EC2 was used to run the backend server of the application, providing scalable computing capacity in the cloud.

Alternatives Considered:

- **AWS Elastic Beanstalk:** Provides a managed service for deploying and scaling web applications, but offers less control over the underlying infrastructure.
- **AWS Fargate:** Manages serverless containers, simplifying infrastructure management, but was less suitable for running the specific backend processes required by **SKEX**.

Why Chosen:

- **Scalability:** EC2 allows for easy scaling up or down based on demand.
- **Flexibility:** Wide range of instance types to choose from based on performance needs.

- **Integration:** Seamless integration with AWS VPC, RDS, and other AWS services ensured efficient deployment and management.

3. AWS Lambda Functions

Description:

AWS Lambda was used to handle serverless functions for various background tasks, such as processing stock data and sending notifications.

Alternatives Considered:

- **Amazon EC2:** Offers more control over the computing environment but requires management of servers.
- **AWS Batch:** Manages batch computing jobs, providing scheduling and resource allocation, but is more suited for large-scale batch processing.

Why Chosen:

- **Cost-Effectiveness:** Pay-as-you-go pricing model reduced costs for infrequent background tasks.
- **Ease of Use:** Simplified the management of code execution without provisioning servers.
- **Integration:** Easily integrates with AWS SNS, RDS, and other services used in the application.

4. Amazon RDS (Relational Database Service)

Description:

Amazon RDS with MariaDB was used to manage the application's database, providing a scalable and reliable relational database service.

Alternatives Considered:

- **Amazon Aurora:** Provides higher performance and availability but at a higher cost, which was not necessary for the initial requirements of SKEX.
- **Amazon DynamoDB:** A fully managed NoSQL database service, which is highly scalable but not suitable for the relational data model used in SKEX.

Why Chosen:

- **Managed Service:** Reduced administrative tasks such as backups, patching, and scaling.
- **Reliability:** High availability and automated backups ensured data integrity and availability.
- **Compatibility:** MariaDB was chosen for its compatibility and performance, fitting well with the application's needs.

5. AWS SNS (Simple Notification Service)

Description:

AWS SNS was used to send notifications to admin about stock transactions and important updates.

Alternatives Considered:

- **Amazon SQS (Simple Queue Service):** Provides message queuing but does not directly handle the publishing and subscribing model needed for notifications.
- **AWS SES (Simple Email Service):** Focuses on email sending capabilities, which could be part of the notification system but lacks the broader messaging capabilities of SNS.

Why Chosen:

- **Ease of Use:** Simple to set up and manage notifications.
- **Integration:** Seamless integration with AWS Lambda and other services.
- **Scalability:** Automatically scales to handle a large number of messages.

6. AWS Backup

Description:

AWS Backup was used to automate and centrally manage backups across AWS services, ensuring data protection and recovery.

Alternatives Considered:

- **Amazon S3:** Provides scalable object storage, but requires custom implementation for backup scheduling and management.
- **AWS Data Lifecycle Manager:** Manages EBS snapshots, providing automated backup management, but does not offer the same centralized management across multiple AWS services as AWS Backup.

Why Chosen:

- **Centralized Management:** Allows for centralized management of backups across multiple AWS services.
- **Automation:** Simplifies the process of scheduling and managing backups.
- **Reliability:** Ensures data protection and easy recovery options.

Deployment Model

The deployment model chosen for the SKEX Trading Application involves a hybrid cloud infrastructure that leverages both public and private resources within AWS. This model provides the flexibility, scalability, and security required for a stock trading platform. The architecture consists of a Virtual Private Cloud (VPC) with public and private subnets, an EC2 instance, Lambda functions, an RDS database, SNS for notifications, and AWS Backup for data protection.

Components

1. AWS VPC

- **Reasoning:**
 - **Isolation and Security:** The VPC provides a logically isolated network that ensures security and control over the network configuration.
 - **Customizability:** Ability to define multiple subnets, route tables, and network gateways to segregate public and private resources.
- **Configuration:**
 - CIDR Block: 10.0.0.0/16
 - Subnets: 3 public and 3 private subnets across different availability zones.

2. EC2 Instance

- **Reasoning:**
 - **Compute Resources:** Required for hosting the application frontend or backend components that need a persistent runtime environment.
 - **Flexibility:** Easy to scale vertically by changing instance types as demand changes.
- **Configuration:**
 - Instance Type: t2.micro (for initial deployment)
 - Security Group: Allows traffic on required ports (e.g., 3000, 5000) and restricts access to only necessary IP ranges.

3. Lambda Functions

- **Reasoning:**
 - **Serverless:** Reduces the need for managing infrastructure, with automatic scaling based on the load.

- **Cost Efficiency:** Charges only for the compute time used, making it cost-effective for event-driven operations.
- **Functions:**
 - buy-stock-test: Handles stock purchase requests.
 - sell-stock-test: Handles stock sell requests.
- **Integration:** Connected to API Gateway to expose these functions as RESTful APIs.

4. RDS Database

- **Reasoning:**
 - **Managed Service:** Simplifies database management tasks such as backups, patching, and scaling.
 - **Reliability:** Provides high availability and durability with multi-AZ deployments.
- **Configuration:**
 - Engine: MariaDB
 - Instance Class: db.r5.large (memory optimized)
 - Security: Deployed in private subnets with restricted access through security groups.

5. AWS SNS (Simple Notification Service)

- **Reasoning:**
 - **Real-time Notifications:** Facilitates real-time notifications to users about significant events like stock transactions.
 - **Scalability:** Can handle a high throughput of messages, ensuring reliable delivery.
- **Configuration:**
 - Topic: skex-sns
 - Subscription: Email notification to dibhavsar214@gmail.com

6. AWS Backup

- **Reasoning:**
 - **Automated Backups:** Ensures regular, automated backups of RDS databases, enhancing data protection and recovery.

- **Compliance:** Helps meet compliance requirements by maintaining backups as per defined policies.
- **Configuration:**
 - Backup Vault: skex-backup-vault
 - Backup Plan: skex-backup-plan with daily backups retained for 30 days.

Deployment Model

The deployment model follows a multi-tier architecture with a separation of concerns between different layers:

1. Network Layer (VPC and Subnets)

- Public subnets host resources that need direct internet access, such as the EC2 instance for web applications.
- Private subnets host the RDS database and other internal services, enhancing security by restricting direct internet access.

2. Compute Layer (EC2 and Lambda)

- The EC2 instance provides a traditional compute environment for hosting persistent applications.
- Lambda functions handle specific, event-driven tasks like stock transactions, ensuring efficient resource utilization.

3. Database Layer (RDS)

- A managed database service (MariaDB) provides reliable and scalable data storage.

4. Notification and Backup Layer (SNS and AWS Backup)

- SNS facilitates user notifications, ensuring timely communication of critical events.
- AWS Backup automates data protection, ensuring regular backups and quick recovery options.

Benefits of Chosen Deployment Model

1. **Scalability:** Both EC2 instances and Lambda functions can be scaled to meet demand, ensuring the system can handle varying loads efficiently.
2. **Security:** The VPC ensures isolation, and security groups restrict access, protecting resources from unauthorized access.

3. **Cost Efficiency:** The combination of on-demand EC2 instances and serverless Lambda functions optimizes costs by paying only for what is used.
4. **Manageability:** AWS managed services like RDS and Backup reduce the operational overhead, allowing the focus to remain on application development and improvement.
5. **Flexibility:** The architecture can be easily adapted to include more services or scale existing ones as the application grows.

This deployment model provides a robust, scalable, and secure foundation for the SKEX Trading Application, ensuring optimal performance and reliability for users.

Delivery Model

The chosen delivery model for the SKEX Trading Application relies on **Docker-based deployment** combined with **Infrastructure as Code (IaaS)** using **AWS CloudFormation**. This model ensures that the application is consistently deployed in a containerized environment, providing portability, scalability, and ease of management.

Components of the Delivery Model

1. Docker-based Deployment

- **Reasoning:**
 - **Portability:** Docker containers encapsulate all the dependencies and environment configurations, ensuring the application runs consistently across different environments.
 - **Isolation:** Each service runs in its own container, providing isolation and reducing the risk of conflicts between services.
 - **Scalability:** Containers can be easily scaled up or down based on demand.
 - **Efficiency:** Docker images can be pulled from a repository, making deployments faster and more efficient.
- **Tools Used:**
 - **Docker:** To create, manage, and run containers.
 - **Docker Hub:** To store and distribute Docker images.

2. Infrastructure as Code (IaC) using AWS CloudFormation

- **Reasoning:**
 - **Reproducibility:** CloudFormation templates ensure consistent and repeatable deployments of the entire infrastructure.
 - **Version Control:** Infrastructure code can be versioned and managed similarly to application code.
- **Tools Used:**
 - **AWS CloudFormation:** To define and provision the AWS infrastructure including VPCs, subnets, EC2 instances, RDS, Lambda functions, SNS topics, and backup configurations.

Docker-based Deployment Workflow

1. Build Docker Images

- Docker images for the application are built and tested locally.

- Images are tagged and pushed to Docker Hub .

2. Pull Docker Images

- During deployment, Docker images are pulled from the repository into the EC2 instances.
- EC2 instances run the Docker containers using the pulled images.

Infrastructure Provisioning Workflow using AWS CloudFormation

1. Define Infrastructure as Code

- CloudFormation templates define the entire infrastructure required for the SKEX application, including:
 - VPC with public and private subnets.
 - Security groups for EC2 instances and RDS.
 - EC2 instances configured to pull and run Docker containers.
 - RDS instance for the database.
 - Lambda functions for specific operations.
 - SNS topics for notifications.
 - Backup configurations for data protection.

2. Deploy Infrastructure

- CloudFormation stacks are deployed to provision the defined infrastructure.
- The stack manages the lifecycle of the infrastructure, allowing updates, deletions, and monitoring.

Benefits of the Chosen Delivery Model

- **Consistency:** Ensures that the application and infrastructure are consistently deployed across different environments.
- **Efficiency:** Reduces the time and effort required for deployment and configuration, allowing developers to focus on application development.
- **Scalability:** Easily scales the application and infrastructure based on demand.
- **Portability:** Docker containers provide a portable environment, making it easier to move the application across different environments or cloud providers.
- **Manageability:** Infrastructure as Code with CloudFormation simplifies the management and monitoring of the AWS resources.

Security Analysis

Security is a critical component of the SKEX stock trading application. Ensuring data security through all stages of the system—whether in transit or at rest—was a key consideration during the development process. This analysis outlines the security measures implemented to protect data and ensure compliance with industry standards.

Data Security in Transit

HTTP Encryption

All data transmitted between the client and the server is secured using HTTP. HTTP ensures that the data is encrypted using Transport Layer Security (TLS), preventing attacks. Key measures include:

API Gateway

For the serverless components, AWS API Gateway is used to expose APIs securely. Key features include:

- **Request Validation:** Validates incoming requests to ensure they meet the defined API specifications.
- **Rate Limiting:** Protects against denial-of-service (DoS) attacks by limiting the number of requests a client can make within a specified time period.
- **Usage Plans and API Keys:** Provides controlled access to the APIs, ensuring only authorized users can interact with the backend services.

Data Security at Rest

Amazon RDS

The application's database is hosted on Amazon RDS, which provides robust security features for data at rest.

Amazon S3 and AWS Backup

For backups and other stored data, Amazon S3 is used along with AWS Backup to ensure data security.

EC2 Instances

The EC2 instances hosting the backend server also follow strict security protocols for data at rest.

IAM Roles: EC2 instances use IAM roles to securely access other AWS services without the need for storing credentials on the instances.

Network Security

AWS VPC

AWS VPC provides a secure network environment for all components of the application. Key measures include:

- **Subnets:** Public and private subnets are used to isolate resources based on their exposure to the internet.
- **Security Groups:** Act as virtual firewalls to control inbound and outbound traffic to the instances.

Authentication and Authorization

AWS IAM

AWS Identity and Access Management (IAM) is used to manage access to AWS services and resources. Key measures include:

- **Fine-Grained Permissions:** IAM policies are used to grant the least privilege necessary to users and services.
- **Multi-Factor Authentication (MFA):** MFA is enabled for administrative access to the AWS management console, ensuring an additional layer of security.

Application-Level Security

The application itself implements robust authentication and authorization mechanisms. Key measures include:

- **User Authentication:** Users sign up with an email, SIN number, and password. Passwords are hashed and stored securely.
- **Session Management:** Secure session management ensures that user sessions are handled securely, preventing session hijacking.

Monitoring and Incident Response

CloudWatch

To ensure continuous monitoring and quick incident response, CloudWatch are utilized. Key measures include:

- **Logging and Monitoring:** All API calls, changes to resources, and operational logs are monitored and recorded.

Cost Analysis

1. AWS VPC

Upfront Costs:

- **Creation of VPC and subnets:** No upfront cost; included in the ongoing costs.

Ongoing Costs:

- **VPC:** No direct cost for the VPC itself.
- **Subnets:** No direct cost for subnets.
- **Internet Gateway:** No direct cost.
- **NAT Gateway:** Not mentioned in the template but if used, costs \$0.045 per hour and \$0.045 per GB data processed.
- **Data Transfer:** Data transfer out to the internet costs \$0.09 per GB.

2. AWS EC2

Upfront Costs:

- **EC2 Instance:** No upfront cost if using On-Demand instances. Reserved instances have upfront payment options.

Ongoing Costs:

- **EC2 Instance (t2.micro):**
 - On-Demand Pricing: Approximately \$0.0116 per hour.
 - Monthly Cost: $\$0.0116 * 24 * 30 = \8.35 (approx.)
- **EBS Storage:**
 - Assuming 30 GB of General Purpose SSD (gp2): \$0.10 per GB-month.
 - Monthly Cost: $30 * \$0.10 = \3.00
- **Data Transfer:**
 - Data transfer out to the internet costs \$0.09 per GB after the first 1 GB per month.

3. AWS Lambda Functions

Upfront Costs:

- **Lambda Function Creation:** No upfront cost.

Ongoing Costs:

- **Lambda Execution:**
 - Memory: 128 MB
 - Duration: 100 ms
 - Number of Requests: Assume 1 million requests per month
- **Cost Calculation:**
 - Compute Charges: 1M requests * 128 MB * 100 ms = 12.8 GB-seconds.
 - Monthly Cost: First 1M requests and 400,000 GB-seconds are free. Beyond that, \$0.20 per 1M requests and \$0.00001667 per GB-second.

Since the example usage falls within the free tier, the monthly cost is \$0.

4. AWS RDS

Upfront Costs:

- **RDS Instance Creation:** No upfront cost if using On-Demand instances. Reserved instances have upfront payment options.

Ongoing Costs:

- **RDS Instance (db.r5.large):**
 - On-Demand Pricing: Approximately \$0.283 per hour.
 - Monthly Cost: $\$0.283 * 24 * 30 = \203.76 (approx.)
- **Storage:**
 - Allocated Storage: 20 GB
 - Cost: \$0.115 per GB-month.
 - Monthly Cost: $20 * \$0.115 = \2.30
- **Backup Storage:**
 - 100% of the allocated storage for free; beyond that, \$0.095 per GB-month.

5. AWS SNS

Upfront Costs:

- **SNS Topic Creation:** No upfront cost.

Ongoing Costs:

- **SNS Usage:**
 - Number of messages published: Assume 1,000 per month
 - Cost: First 1 million requests per month are free. Beyond that, \$0.50 per 1M requests.
 - Email Delivery: \$0.0004 per email
- **Monthly Cost:** Assuming 1,000 emails.
 - SNS Requests: Free
 - Email Cost: $1,000 * \$0.0004 = \0.40

6. AWS Backups

Upfront Costs:

- **Backup Plan Creation:** No upfront cost.

Ongoing Costs:

- **Backup Storage:**
 - Assume 20 GB of RDS snapshot storage
 - Cost: \$0.095 per GB-month.
 - Monthly Cost: $20 * \$0.095 = \1.90

Total Monthly Costs:

- **VPC:** Minimal costs, mostly from data transfer.
- **EC2 Instance:** $\$8.35 + \$3.00 = \$11.35$
- **Lambda Functions:** \$0 (within free tier)
- **RDS Instance:** $\$203.76 + \$2.30 = \$206.06$
- **SNS:** \$0.40
- **Backups:** \$1.90

Total Monthly Estimate: $\$11.35 + \$206.06 + \$0.40 + \$1.90 = \$219.71$ (approx.)

Alternative Approaches

1. **EC2 Reserved Instances:**
 - Switching to Reserved Instances for EC2 could save up to 75% over On-Demand prices, depending on the term and payment options.

2. Using Smaller RDS Instances:

- Downgrading the RDS instance to a smaller class (e.g., db.t3.medium) could significantly reduce costs if the performance requirements are not as high.

3. Lambda Cost Optimization:

- Ensuring Lambda functions are optimized to run in the least amount of time and with the minimum necessary memory can help reduce costs further.

4. Use of Spot Instances for EC2:

- Spot Instances can be used for non-critical workloads, potentially reducing EC2 costs by up to 90%.

5. Consolidating SNS and Lambda Functions:

- Combining multiple related functionalities into fewer Lambda functions and minimizing SNS usage could further optimize costs.

Justification for the Solution

1. Scalability and Flexibility:

- Using On-Demand and t2.micro instances ensures flexibility and scalability for fluctuating workloads.

2. High Availability and Performance:

- The choice of db.r5.large for RDS ensures high performance and reliability, which is crucial for database operations.

3. Compliance and Data Protection:

- Implementing AWS Backups with regular snapshots ensures data protection and compliance with best practices for disaster recovery.